

A Fast Seeded Graph Matching Algorithm, and Seeded Matching of Random Graphs from a Common Distribution

Donniell's writeup, in progress

July 22, 2012

Abstract

Given two graphs on the same number of vertices, the graph matching problem is to find a bijection between the two vertex sets which minimizes the number of adjacency disagreements between the two graphs; this problem is NP-hard. The seeded graph matching problem is the graph matching problem with an additional constraint that the bijection assigns some particular vertices of one vertex set to respective particular vertices of the other vertex set. We modify the state-of-the-art approximate graph matching algorithm of Vogelstein, Conroy, et al to make it an approximate seeded graph matching algorithm.

Then we illustrate that two random graphs drawn from the same distribution might be graph matched to each other in a way that almost completely differs from the natural underlying correspondence between the two graph's vertices within the random graph model. However, seeding some of the vertices according to this natural, underlying correspondence can very significantly reduce—for the nonseeded vertices—the difference between the seeded graph matching and the underlying correspondence. The effectiveness of the modified Vogelstein, Conroy et al approximate seeded graph matching algorithm in this context is illustrated with simulations, and through an example of matching a library of linked Wikipedia documents in English to the corresponding linked Wikipedia documents in French, an example of matching an electrical network to its chemical counterpart in a *Caenorhabditis elegans* (a particular roundworm) brain, and an example of fraud detection by matching communication graphs from the failed energy company Enron.

1 Setting and Overview

All graphs in this manuscript are simple, ie the edges are undirected and there are no loops or multiple edges. In other words, the adjacency matrices for graphs are $\{0, 1\}$ -valued square matrices that are symmetric, with only zeros along the main diagonal. This restriction to simple graphs is just a convenience; indeed, all of our work and analysis can be naturally extended to settings with more general graphs.

Suppose G_1 and G_2 are two graphs with respective vertex sets V_1 and V_2 such that $\#V_1 = \#V_2$. For any bijective function $\phi : V_1 \rightarrow V_2$, the number of adjacency disagreements under ϕ is defined to be $d(\phi) := \#\{(u, v) \in V_1 \times V_1 : [u \sim_{G_1} v \text{ and } \phi(u) \not\sim_{G_2} \phi(v)] \text{ or } [u \not\sim_{G_1} v \text{ and } \phi(u) \sim_{G_2} \phi(v)]\}$. The *graph matching problem* is to minimize $d(\phi)$ over all bijective functions $\phi : V_1 \rightarrow V_2$. This problem is NP-hard; in fact, even the weaker problem of just deciding whether there exists a graph isomorphism between G_1 and G_2 is notoriously of unknown complexity (and, indeed, is suspected to belong to an intermediate complexity class which is strictly between P and NP-complete). Thus, in particular, there are no efficient algorithms known for graph matching, and it is not expected that any exist.

Developing graph matching heuristics is a burgeoning field. An excellent survey article by Conte, Foggia, Sansone, and Vento titled “Thirty years of graph matching in pattern recognition” [2] outlines successful application of approximate graph matching to two-dimensional and three-dimensional image analysis, document processing, biometric identification, image databases, video analysis, and biological and biomedical applications. The current state-of-the-art algorithms can provide effective and realtime approximate graph matching for graphs on $\approx 10^3$ vertices.[3]

In this manuscript, we will utilize the approximate graph matching algorithm of Vogelstein, Conroy et al [3] which they call “FAQ” (an acronym for Fast Approximate Quadratic Assignment Problem Algorithm); its running time is cubic in the number of vertices and, in practice, the quality of the approximate solution and the speed of the algorithm are state-of-the-art. The details of FAQ will be specified later, in Section 4.

Suppose we are also given subsets $W_1 \subset V_1$, $W_2 \subset V_2$ such that $\#W_1 = \#W_2$, and we are given a fixed bijection $\psi : W_1 \rightarrow W_2$. We define the *seeded graph matching problem* as the problem of minimizing $d(\phi)$ over all bijections $\phi : V_1 \rightarrow V_2$ that agree with ψ on W_1 (ie, for all $u \in W_1$, $\phi(u) = \psi(u)$). The elements of W_1 are called *seeds* and the bijection ψ is a *seeding*. In Section 4, we modify the approximate graph matching FAQ algorithm for use in approximate seeded graph matching.

When we say that “ G_1 on vertex set V_1 , and G_2 on vertex set V_2 are random graphs independently drawn from the same distribution, with *correspondence function* Ψ ”, (for the bijective function $\Psi : V_1 \rightarrow V_2$), we mean that there are specified probabilities for each of the $2^{\binom{\#V_1}{2}}$ possible

graphs on the vertex set V_1 and, from this probability distribution, the two graphs G_1 and G_2 are independently realized and then—just in G_2 —each vertex $u \in V_1$ is relabeled as $\Psi(u) \in V_2$, so that G_1 remains on vertex set V_1 but G_2 now has vertex set V_2 . The approximate graph matching solution $\phi : V_1 \rightarrow V_2$ may be viewed as an approximation for the underlying correspondence function $\Psi : V_1 \rightarrow V_2$, if Ψ is partially or completely unknown.

We will see in Section 5 that ϕ may be a very lousy approximation for Ψ in general, perhaps agreeing with Ψ at only a handful of vertices, not much different from chance. However, we will also see that utilizing some seeds $W_1 \subset V_1$ —and the seeding function ψ consisting of the restriction of Ψ to W_1 —can yield an approximate seeded graph match solution which agrees with Ψ on a much more substantial fraction of the nonseeded vertices from V_1 .

The structure of this paper is as follows:

In Section 2 we formulate the relaxation of the seeded graph matching problem. In Section 3 we describe a straight linear programming approach—which we call “SLP”—to approximately solve the seeded graph matching problem. SLP is guaranteed to optimally solve relaxed graph matching problems (in contrast to all current state-of-the-art algorithms), and SLP runs in polynomial time. Thus, assuming that we give up on the intractable unrelaxed (integral) seeded graph matching problem in favor of solving a relaxation and projecting it to the nearest integer solution, then it can be argued that SLP is the best choice for this *if we don’t care at all about running time except to demand that running time be polynomial*. Thus, SLP will serve as an excellent “straw man” here, a benchmark to compare other algorithms against to assess quality of solution (even though SLP is many orders of magnitude slower than state-of-the-art).

In Section 4 we adapt the FAQ algorithm of Vogelstein, Conroy et al into an algorithm for approximate seeded graph matching; in Section 5 we observe that modified FAQ seems to be just as effective as SLP (while running many orders of magnitude faster). We see that for random graphs drawn from the same distribution, the approximate graph matching will not in general agree much with the underlying correspondence function, but seeding can dramatically improve such agreement for the nonseeded vertices. Examples that illustrate this are (Section 6) matching a library of linked Wikipedia documents in English to the corresponding linked Wikipedia documents in French, (Section 7) an example of matching an electrical network to its chemical counterpart in a *Caenorhabditis elegans* brain, and (Section 8) an example of fraud detection by matching communications networks from the failed energy company Enron.

2 The relaxed seeded graph matching problem

We are interested in solving the seeded graph matching problem but, as discussed earlier, the seeded graph matching problem is NP-hard, and we thus have no expectation that there even exists an efficient algorithm. So we seek approximate solutions that can be efficiently computed.

In this section we express the seeded graph matching problem as an optimization problem with integer constraints, and then we relax the integer constraints by replacing them with nonnegativity constraints. Of course, when solving a relaxation, the solution may not in general be integer valued, and as such it would not even be appropriate as an approximate solution to the unrelaxed original problem. However, we will then project the solution of the relaxed optimization problem to the nearest member of the feasible region of the unrelaxed problem, and we then declare that to be the approximate solution of the original (unrelaxed) problem.

Recall that G_1 is a graph on vertex set V_1 , G_2 is a graph on vertex set V_2 such that $\#V_1 = \#V_2$, the set of seeds W_1 is a subset of V_1 , W_2 is a subset of V_2 such that $\#W_1 = \#W_2$, and bijection $\psi : W_1 \rightarrow W_2$ is the seeding. Without loss of generality, we will take V_1 and V_2 to each be the set of integers $\{1, 2, \dots, m+n\}$, we will take W_1 and W_2 to each be the set of integers $\{1, 2, \dots, m\}$, we will take ψ to be the identity function, for some fixed nonnegative integer m and positive integer n . Let $A, B \in \mathbb{R}^{(m+n) \times (m+n)}$ be the adjacency matrices for G_1 and G_2 , respectively; this means that for all $i, j \in \{1, 2, \dots, m+n\}$ it holds that $a_{ij} = 1$ or 0 according as $i \sim_{G_1} j$ or not, and $b_{ij} = 1$ or 0 according as $i \sim_{G_2} j$ or not. It will soon be useful to let A and B be partitioned as

$$A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \quad B = \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

where $A_{11}, B_{11} \in \mathbb{R}^{m \times m}$, $A_{12}, B_{12} \in \mathbb{R}^{m \times n}$, $A_{21}, B_{21} \in \mathbb{R}^{n \times m}$, and $A_{22}, B_{22} \in \mathbb{R}^{n \times n}$.

It is clear that the seeded graphmatch problem is to minimize $\|A - (I_{m \times m} \oplus P)B(I_{m \times m} \oplus P)^T\|_1$ over all $n \times n$ permutation matrices P , where $I_{m \times m}$ is the m -by- m identity matrix, \oplus is the direct sum of matrices, and $\|\cdot\|_1$ is the ℓ_1 vector norm on matrices; say the optimal P is \tilde{P} . Then the corresponding bijection $\phi_{\tilde{P}} : \{1, 2, \dots, m+n\} \rightarrow \{1, 2, \dots, m+n\}$ defined as, for all $i \in \{1, 2, \dots, m\}$, $\phi_{\tilde{P}}(i) = i$ and, for all $i, j \in \{1, 2, \dots, n\}$, $\phi_{\tilde{P}}(i+m) = j+m$ precisely when $\tilde{p}_{ij} = 1$, is the bijection which solves the seeded graphmatch problem.

Of course, this optimization problem is equivalent to minimizing $\|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_1$ or $\|A - (I_{m \times m} \oplus P)B(I_{m \times m} \oplus P)^T\|_2$ or $\|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_2$, over all permutation matrices P , where $\|\cdot\|_2$ is the ℓ_2 vector norm on matrices. Expanding out $\|A - (I_{m \times m} \oplus P)B(I_{m \times m} \oplus P)^T\|_2^2 = \|A\|_2^2 + \|B\|_2^2 - 2 \cdot \text{trace} A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$, we see that this optimization problem is equivalent to maximizing $\text{trace} A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$ over permutation matrices P .¹

¹Note that although A and B are symmetric matrices, we nonetheless keep transposes in place wherever they

Although we don't expect to ever find an efficient algorithm for seeded graph matching, we will next have efficient algorithms for solving a relaxation of seeded graph matching. In Section 3, where we define the algorithm SLP, we will minimize $\|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_1$ subject to the constraint that P is a doubly stochastic matrix, which means that $P \in \mathbb{R}^n$ such that $P\vec{1}_n = \vec{1}_n$, $P^T\vec{1}_n = \vec{1}_n$, and $P \geq 0_{n \times n}$ coordinatewise, where $0_{n \times n}$ is the n -by- n matrix of zeros and $\vec{1}_n$ is the n -vector of all ones. In Section 4, where we define the modified FAQ algorithm, we maximize $\text{trace}A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$ over all doubly stochastic matrices P . Indeed, both of these are relaxations of the seeded graphmatch problem in the sense that if we were to add integrality constraints—that P is integer-valued—then we would precisely return to the constraint that P is a permutation matrix, hence we would have returned to the seeded graphmatch problem.

There is one last detail to cover in this section. We will indeed efficiently solve these relaxations of the seeded graph matching problem, say the solution is the doubly stochastic matrix \tilde{P} ; but, if \tilde{P} is not a permutation matrix, then how do we get out of \tilde{P} a meaningful approximate solution to the seeded graph matching problem? The answer is that we will do one more step; we will find the permutation matrix \tilde{Q} which solves the optimization problem $\min \|Q - \tilde{P}\|_1$ subject to Q being a permutation matrix, and finally $\phi_{\tilde{Q}}$ is our approximate seeded graphmatch solution. To solve this latter optimization problem, observe that for any permutation matrix Q

$$\begin{aligned}
\|Q - \tilde{P}\|_1 &= \sum_{i,j \in \{1,2,\dots,n\}: q_{ij} \neq 1} \tilde{p}_{ij} + \sum_{i,j \in \{1,2,\dots,n\}: q_{ij} = 1} (1 - \tilde{p}_{ij}) \\
&= \sum_{i,j \in \{1,2,\dots,n\}} \tilde{p}_{ij} + \sum_{i,j \in \{1,2,\dots,n\}: q_{ij} = 1} (1 - 2\tilde{p}_{ij}) \\
&= n + n - 2 \cdot \sum_{i,j \in \{1,2,\dots,n\}: q_{ij} = 1} \tilde{p}_{ij} \\
&= 2n - 2 \text{trace}Q^T \tilde{P}.
\end{aligned}$$

Thus, minimizing $\|Q - \tilde{P}\|_1$ subject to Q being a permutation matrix is equivalent to maximizing $\text{trace}Q^T \tilde{P}$ subject to Q being a permutation matrix; this latter optimization formulation is precisely a formulation of the well-known linear assignment problem, and it is efficiently solvable in $O(n^3)$ time with Edmonds and Karp's [cite] and Tomizawa's [cite] implementation of the so-called Hungarian Algorithm. In this manner we can efficiently obtain $\phi_{\tilde{Q}}$, which is our approximate seeded graphmatch solution.

are present to enable further generalization; our analysis will not change if we instead were in a broader setting where A and B are generic (nonsymmetric, nonhollow, and/or nonintegral) matrices in $\mathbb{R}^{(m+n) \times (m+n)}$.

3 SLP algorithm: the straw man

This section describes the straight linear programming “SLP” approach for approximate seeded graph matching. In SLP, we first optimally solve the relaxed seeded graph matching problem: minimize $\|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_1$ subject to the constraint that P is a doubly stochastic matrix. This is done by solving the linear programming problem described below with an interior point method. Then SLP projects the linear program’s solution to the nearest permutation matrix and converts that permutation matrix to a seeded graph matching ϕ in the manner described at the end of Section 2. This ϕ is the output of SLP. (Straight linear programming has historically been used to tackle (nonseeded) graph matching in general, see eg [1].)

SLP runs in polynomial time since it just involves solving one linear programming problem with an interior point method and then one use of the Hungarian Algorithm to project the linear program’s solution to a permutation matrix. However, solving this linear program directly is relatively expensive, and therefore the running time of SLP is not even remotely as fast as the modified FAQ algorithm described in Section 4. However, SLP is guaranteed to optimally solve the relaxed seeded graph matching problem in polynomial time (as opposed to all of the other state-of-the-art algorithms, which may not converge or may converge to stationary points or local minima that are not globally optimal) and, for this reason, we use SLP as the “straw man” to compare the quality of the output of the modified FAQ algorithm against the quality of the output of SLP. Indeed, if we did not care about running time except that it be polynomial, then an argument can be made that SLP is the right algorithm to use for approximate seeded graph matching.

We now spell out the linear program that expresses minimize $\|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_1$ subject to the constraint that P is a doubly stochastic matrix. Note that

$$\begin{aligned} \|A(I_{m \times m} \oplus P) - (I_{m \times m} \oplus P)B\|_1 &= \left\| \begin{bmatrix} A_{11} & A_{12}P \\ A_{21} & A_{22}P \end{bmatrix} - \begin{bmatrix} B_{11} & B_{12} \\ PB_{21} & PB_{22} \end{bmatrix} \right\|_1 \\ &= \|A_{11} - B_{11}\|_1 + \|(I_{n \times n} \otimes A_{22} - B_{22}^T \otimes I_{n \times n}) \text{Vec}P\|_1 \\ &\quad + \|((I_{n \times n} \otimes A_{12})\text{Vec}P) - \text{Vec}B_{12}\|_1 + \|((B_{21}^T \otimes I_{n \times n})\text{Vec}P) - \text{Vec}A_{21}\|_1 \end{aligned}$$

where \otimes denotes the Kronecker product of two matrices, and “Vec” of a matrix denotes the concatenation the columns of the argument matrix into a single, long column. By introducing two sets of slack variables ϵ^+ and ϵ^- to capture ℓ_1 discrepancy, and by noting that the conditions $P\vec{1}_n = \vec{1}_n$ and $P^T\vec{1}_n = \vec{1}_n$ can be expressed as $(I_{n \times n} \otimes \vec{1}_n^T)\text{Vec}P = \vec{1}_n$ and $(\vec{1}_n^T \otimes I_{n \times n})\text{Vec}P = \vec{1}_n$, we thus obtain relaxed seeded graphmatch as the following linear program in standard form:

$$\begin{aligned}
& \min \begin{bmatrix} \vec{0}_{n^2} \\ \vec{1}_{n^2+2mn} \\ \vec{1}_{n^2+2mn} \end{bmatrix}^T \begin{bmatrix} \text{vec} P \\ \epsilon^+ \\ \epsilon^- \end{bmatrix} \\
& \text{such that} \begin{bmatrix} I_{n \times n} \otimes A_{22} - B_{22}^T \otimes I_{n \times n} & & & & \\ & I_{n \times n} \otimes A_{12} & & & \\ & B_{21}^T \otimes I_{n \times n} & & & \\ & I_{n \times n} \otimes \vec{1}_n^T & & & \\ & \vec{1}_n^T \otimes I_{n \times n} & & & \end{bmatrix} \begin{bmatrix} I_{(n^2+2mn) \times (n^2+2mn)} & -I_{(n^2+2mn) \times (n^2+2mn)} \\ & 0_{2n \times (n^2+2mn)} & 0_{2n \times (n^2+2mn)} \end{bmatrix} \begin{bmatrix} \text{vec} P \\ \epsilon^+ \\ \epsilon^- \end{bmatrix} = \begin{bmatrix} \vec{0}_{n^2} \\ \text{vec} B_{12} \\ \text{vec} A_{21} \\ \vec{1}_n \\ \vec{1}_n \end{bmatrix} \\
& \text{and} \begin{bmatrix} \text{vec} P \\ \epsilon^+ \\ \epsilon^- \end{bmatrix} \geq \vec{0}_{3n^2+4mn} \quad \text{where } \text{vec} P \in \mathbb{R}^{n^2} \text{ and } \epsilon^+, \epsilon^- \in \mathbb{R}^{n^2+2mn}.
\end{aligned}$$

This linear program can be efficiently solved by an interior point method. This concludes the description of SLP.

4 Modified-FAQ for seeded graph matching

In this section we modify the state-of-the-art approximate graph matching algorithm of Vogelstein, Conroy, et al, which they call FAQ, so that it can be used for approximate seeded graph matching. As mentioned before, the seeded graph matching problem can be expressed as that of maximizing $\text{trace} A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$ over all permutation matrices P and, as mentioned before, we don't expect that there is an efficient algorithm for this, so we instead consider the relaxation, which is maximizing $\text{trace} A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$ over all doubly stochastic matrices P . When a solution (or approximate solution) to the relaxed problem is found, then it is projected to the nearest permutation matrix and converted to a seeded graph matching ϕ in the manner described at the end of Section 2. This ϕ is the output of modified-FAQ.

The modified FAQ algorithm solves (or approximately solves) the relaxed seeded graph matching problem—maximize $\text{trace} A^T(I_{m \times m} \oplus P)B(I_{m \times m} \oplus P^T)$ subject to P being a doubly stochastic matrix—by using the Frank-Wolfe Method, which is an iterative procedure that involves successively solving linearizations. It turns out that the linearizations can be cast as linear assignment problems that are solved with the Hungarian Algorithm. So, although SLP involves solving only one linear program and here we will solve a succession of linear programs, nonetheless the special form of the linear programs in this section allows for the Hungarian Algorithm to be employed, which is orders of magnitude faster than an interior point method for solving a linear program.

We first briefly review the Frank-Wolfe Method before proceeding to apply it. The general kind of optimization problem for which the Frank-Wolfe Method is used is

$$(FWP) \quad \text{Minimize } f(x) \text{ such that } x \in S, \quad (1)$$

where S is a polyhedral set (ie is described by linear constraints) in a Euclidean space of some dimension, and the function $f : S \rightarrow \mathbb{R}$ is continuously differentiable. A starting point $x^{(1)} \in S$ is chosen in some fashion, perhaps arbitrarily. For $i = 1, 2, 3, \dots$, the following is done. The function $\tilde{f}^{(i)} : S \rightarrow \mathbb{R}$ is defined to be the first order (ie linear) approximation to f at $x^{(i)}$ —that is, $\tilde{f}^{(i)}(x) := f(x^{(i)}) + \nabla f(x^{(i)})^T(x - x^{(i)})$; then solve the linear program: minimize $\tilde{f}^{(i)}(x)$ such that $x \in S$ (this can be done efficiently since it is a linear objective function with linear constraints, and note that, by ignoring additive constants, the objective function of this subproblem can be abbreviated: minimize $\nabla f(x^{(i)})^T x$ such that $x \in S$), say the solution is $\tilde{x}^{(i)} \in S$. Now, the point $x^{(i+1)} \in S$ is defined as the solution to: minimize $f(x)$ such that x is on the line segment from $x^{(i)}$ to $\tilde{x}^{(i)}$ in S . (This is a just a one dimensional optimization problem; in the case where f is quadratic this can be exactly solved analytically.) Go to the next i , and terminated this iterative procedure when the sequence of iterates $x^{(1)}, x^{(2)}, x^{(3)}, \dots$ stops changing much or develops a gradient close enough to zero. This concludes our review of the Frank-Wolfe Method.

We now describe how modified FAQ employs the Frank-Wolfe Method to solve the relaxed seeded graph matching problem. The objective function here is

$$\begin{aligned} f(P) &= \text{trace} \left(\begin{bmatrix} A_{11}^T & A_{21}^T \\ A_{12}^T & A_{22}^T \end{bmatrix} \begin{bmatrix} I_{m \times m} & 0_{m \times n} \\ 0_{n \times m} & P \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} \begin{bmatrix} I_{m \times m} & 0_{m \times n} \\ 0_{n \times m} & P^T \end{bmatrix} \right) \\ &= \text{trace} \left(\begin{bmatrix} A_{11}^T & A_{21}^T \\ A_{12}^T & A_{22}^T \end{bmatrix} \begin{bmatrix} B_{11} & B_{12}P^T \\ PB_{21} & PB_{22}P^T \end{bmatrix} \right) \\ &= \text{trace} A_{11}^T B_{11} + \text{trace} A_{21}^T P B_{21} + \text{trace} A_{12}^T B_{12} P^T + \text{trace} A_{22}^T P B_{22} P^T \\ &= \text{trace} A_{11}^T B_{11} + \text{trace} P^T A_{21} B_{21}^T + \text{trace} P^T A_{12}^T B_{12} + \text{trace} A_{22}^T P B_{22} P^T \end{aligned}$$

which has gradient

$$\nabla(P) := A_{21} B_{21}^T + A_{12}^T B_{12} + A_{22} P B_{22}^T + A_{22}^T P B_{22}.$$

We start the Frank-Wolfe Algorithm at the doubly stochastic matrix $\tilde{P} = \frac{1}{n} \vec{1}_n \vec{1}_n^T$. (This is only for simplicity, and any other choice of doubly stochastic \tilde{P} might be as effective). In the next paragraph we describe a single step in the Frank-Wolfe algorithm. Such steps are repeated iteratively until the iterates empirically converge.

Given any particular doubly stochastic matrix $\tilde{P} \in \mathbb{R}^{n \times n}$ the Frank-Wolfe-step linearization involves maximizing $\text{trace} Q^T \nabla(\tilde{P})$ over all of the doubly stochastic matrices $Q \in \mathbb{R}^{n \times n}$. This is precisely the linear assignment problem (since it is not hard to show that the optimal doubly stochastic

Q can in fact be selected to be a permutation matrix) and so the Hungarian Algorithm will in fact find the optimal Q , call it \tilde{Q} . The next task in the Frank-Wolfe algorithm step will be maximizing the objective function over the line segment from \tilde{P} to \tilde{Q} ; ie maximizing $g(\alpha) := f(\alpha\tilde{P} + (1-\alpha)\tilde{Q})$ over $\alpha \in [0, 1]$. Denote $c := \text{trace}A_{22}^T\tilde{P}B_{22}\tilde{P}^T$ and $d := \text{trace}(A_{22}^T\tilde{P}B_{22}\tilde{Q}^T + A_{22}^T\tilde{Q}B_{22}\tilde{P}^T)$ and $e := \text{trace}A_{22}^T\tilde{Q}B_{22}\tilde{Q}^T$ and $u := \text{trace}(\tilde{P}^TA_{21}B_{21}^T + \tilde{P}^TA_{12}^TB_{12})$ and $v := \text{trace}(\tilde{Q}^TA_{21}B_{21}^T + \tilde{Q}^TA_{12}^TB_{12})$. Then (ignoring the additive constant $\text{trace}A_{11}^TB_{11}$ without loss of generality, since it won't affect the maximization) we have $g(\alpha) = c\alpha^2 + d\alpha(1-\alpha) + e(1-\alpha)^2 + u\alpha + v(1-\alpha)$ which simplifies to $g(\alpha) = (c-d+e)\alpha^2 + (d-2e+u-v)\alpha + (e+v)$. Setting the derivative of g to zero yields potential critical point $\tilde{\alpha} := \frac{-(d-2e+u-v)}{2(c-d+e)}$ (if indeed $0 \leq \tilde{\alpha} \leq 1$); thus the next Frank-Wolfe algorithm iterate will either be \tilde{P} (in which case algorithm would halt) or \tilde{Q} or $\tilde{\alpha}\tilde{P} + (1-\tilde{\alpha})\tilde{Q}$, and the objective functions can be compared to decide which of these three matrices will be the \tilde{P} of the next Frank-Wolfe step.

At the termination of the Frank-Wolfe Algorithm, we have a solution, or an approximate solution, to the relaxed seeded graph matching problem. This concludes the description of the modified FAQ algorithm. By limiting the number of Frank-Wolfe steps to a constant, the running time is cubic in the number of vertices, since that is the complexity of the Hungarian Algorithm. There is no appreciable difference in running time between modified FAQ and FAQ, and we thus have state-of-the-art running time, in practice.

5 Seeding of random graphs

In this section, we illustrate that for random graphs drawn from the same distribution, the approximate graph matching may not in general agree much with the underlying correspondence function, but seeding might dramatically improve such agreement for the nonseeded vertices.

5.1 Bernoulli random graphs

If a matrix $\Omega \in [0, 1]^{k \times k}$, for some positive integer k , is symmetric with zeros along the main diagonal then a *Bernoulli*(Ω) random graph is a random graph on the vertices $\{1, 2, \dots, k\}$ wherein, for all i, j such that $i < j$, the probability that $i \sim j$ is Ω_{ij} , independently for all such pairs i, j .

Given such a matrix of parameters $\Omega \in [0, 1]^{k \times k}$ and a positive integer n less than k , consider the following experiment, which we will call Experiment(Ω, n): Independently realize two Bernoulli(Ω) random graphs. Call the first graph G_1 . Call the second graph G_2 , but only after permuting/relabeling its vertices with an arbitrarily selected bijective correspondence function $\Psi : \{1, 2, \dots, k\} \rightarrow \{1, 2, \dots, k\}$, where Ψ is arbitrary except that for all $i = 1, 2, \dots, k - n$ it holds that $\Psi(i) = i$. Then, for each of $m = 0, 1, 2, \dots, k - n$, compute the modified-FAQ

seeded graph matching $\phi^{(m)}$ between the subgraphs of G_1 and G_2 induced by the $m + n$ vertices $V_1^{(m)} = V_2^{(m)} := \{k - n - m + 1, k - n - m + 2, k - n - m + 3, \dots, k\}$, using the seeding Ψ on the m seed vertices $W_1^{(m)} = W_2^{(m)} := \{k - n - m + 1, k - n - m + 2, k - n - m + 3, \dots, k - n\}$ (and there are n nonseeded vertices $\{k - n + 1, k - n + 2, k - n + 3, \dots, k\}$). The *match ratio* $\delta^{(m)} := \frac{\#\{v \in V_1^{(m)} \setminus W_1^{(m)} : \phi^{(m)}(v) = \Psi(v)\}}{n}$ is the fraction of the n nonseeded vertices where $\phi^{(m)}$ agrees with correspondence function Ψ . Note that there are always the same n nonseeded vertices that we are interested in matching across the two graphs and, as $m = 0, 1, 2, \dots$, each increment of m adds to the utilized instance a new seed vertex, which is seeded to agree with the assignment function Ψ . When $m = 0$ the graph matching problem instance isn't seeded at all.

In Figure XXa, we plotted the match ratio $\delta^{(m)}$ against $m = 0, 1, 2, \dots, 90$, averaged over 65 repetitions of Experiment($\Omega, 30$), where $k = 120$ and each of the lower triangular elements of each of the 65 parameter matrices Ω was independent, Uniformly distributed on the interval $(0, 1)$. In Figure XXb, we did the same as was done to get Figure XXa, except that we used SLP instead of modified-FAQ to do the seeded graph matching. Several observations should be made:

Note that if instead of graph matching we just discrete-uniform-randomly selected any bijection between the unseeded vertices of the two graphs, the expected number of agreements between such a random bijection and the assignment function is exactly 1. Thus, here “chance” would have match ratio $\frac{1}{30} = .0333$, which is not that much different than the match ratio $\delta^{(0)} \approx .05$ when there were no seeds. As seeds are added here, there was a dramatic increase in the match ratio. When there were around 30 seeds then the match ratio here was more than .50, and where there were around 60 seeds then the match ratio here was more than .90, and there empirically seems to be a close-to-linear relationship between $\delta^{(m)}$ and m , for m between 0 and 60.

ADD BRIEF EXPLANATION AS TO WHY THIS IS SURPRISING/ NOT SURPRISING

Also note that, unlike SLP, modified-FAQ is not guaranteed to optimally solve the relaxed seeded graph matching problem but, nonetheless, the empirical performance here of modified-FAQ was not that different than the empirical performance of SLP. However, on a personal computer, the computing time to perform the experiments for Figure XXb with SLP was around 80 hours, whereas the computing time to perform the experiments for Figure XXa with modified-FAQ was only 134 seconds.

5.2 Two-class random graphs

Suppose that we are given a positive integer k , a function $\gamma : \{1, 2, \dots, k\} \rightarrow \{0, 1\}$, and also p, q, r in the real interval $[0, 1]$. A $\text{Two-Class}(\gamma, p, q, r)$ random graph is a random graph on the vertices $\{1, 2, \dots, k\}$ wherein, for all i, j such that $i < j$, the probability that $i \sim j$ is p, q , or r , according as $\gamma(i) = \gamma(j) = 0$, or $\gamma(i) = \gamma(j) = 1$, or $\gamma(i) \neq \gamma(j)$, independently for all such pairs i, j . The vertices in $\{i : 1 \leq i \leq k, \gamma(i) = 0\}$ will be called the *first class* of vertices, and the vertices in $\{i : 1 \leq i \leq k, \gamma(i) = 1\}$ will be called the *second class* of vertices. The graph might represent, for instance, a social network, with the first class consisting of innocent people and the second class consisting of scheming fraudsters, where the probability of two vertices/people communicating just depends on whether they are both nonfraudsters, or both fraudsters, or exactly one of them is a fraudster. (In Section 8 we will deal with an actual social network of email messages/ edges among Enron employees/ vertices, where some of the employees are innocent and some of them were scheming to commit fraud.)

Given such two-class-random-graph parameters γ, p, q, r and given a positive integer n less than k , consider the following experiment, which we call $\text{Experiment}(\gamma, p, q, r, n)$: Independently realize two $\text{Two-Class}(\gamma, p, q, r)$ random graphs. Call the first graph G_1 . Call the second graph G_2 , but only after permuting/relabeling its vertices with an arbitrarily selected bijective correspondence function $\Psi : \{1, 2, \dots, k\} \rightarrow \{1, 2, \dots, k\}$; Ψ is arbitrary except that for all $i = 1, 2, \dots, k - n$ it holds that $\Psi(i) = i$. Then, for each of $m = 0, 1, 2, \dots, k - n$, compute the modified-FAQ seeded graph matching $\phi^{(m)}$ between the subgraphs of G_1 and G_2 induced by the $m + n$ vertices $V_1^{(m)} = V_2^{(m)} := \{k - n - m + 1, k - n - m + 2, k - n - m + 3, \dots, k\}$, using the seeding Ψ on the m seed vertices $W_1^{(m)} = W_2^{(m)} := \{k - n - m + 1, k - n - m + 2, k - n - m + 3, \dots, k - n\}$ (and there are n nonseeded vertices $\{k - n + 1, k - n + 2, k - n + 3, \dots, k\}$). Of course, all first class vertices are stochastically indistinguishable one from another, and all second class vertices are stochastically indistinguishable one from another. Thus, our goal ought not be to have $\phi^{(m)}$ agree with the correspondence function Ψ but, rather, our goal is that they agree about the class of as many vertices as possible. With this in mind, we define the match ratio to be $\delta^{(m)} := \frac{\#\{v \in V_1^{(m)} \setminus W_1^{(m)} : \gamma(v) = 1, \gamma(\Psi^{-1}(\phi^{(m)}(v))) = 1\}}{\#\{v \in V_1^{(m)} \setminus W_1^{(m)} : \gamma(v) = 1\}}$, ie the fraction of the nonseed second-class vertices that are matched to vertices which correspond to (likewise) second-class vertices.

INSERT SIMULATIONS FOR TWO-CLASS RANDOM GRAPHS SIMILAR TO SIMULATIONS FOR BERNOULLI RANDOM GRAPHS.

6 Wikipedia example

Wikipedia is an online editable encyclopedia with 22 million articles (more than 4 million articles in English) in 285 languages. A collection of $k = 1382$ English articles were collected by crawling the (directed) 2-neighborhood of the document “Algebraic Geometry”. This first graph will be made a simple undirected graph by symmetrizing its adjacency matrix. In Wikipedia, links between articles of the same topic in different languages are available. Thus, 1-1 correspondence information between the vertices of these English wikipedia subgraph and some vertices of the French wikipedia graph is available. Corresponding articles in French were collected to form a second graph which is not necessarily connected. Following the notation in previous sections, the English wikipedia subgraph is G_1 and the French wikipedia subgraph induced by the correspondents of the English wikipedia articles is G_2 . $n = 500$ vertices are randomly chosen as nonseeded vertices, and $m = 0, \dots, 882$ seeds are randomly sampled without replacement as seeds. Fig 2 shows the plot of the *match ratio* vs the number of hard seeds.

This example is the most computationally expensive of our examples. Figure ?? shows the running time vs number of vertices to be matched.

7 Caenorhabditis elegans brain example

C. elegans is a worm that is a model organism that has been extensively studied. Its particular usefulness comes from its nervous system, consisting of 302 neurons whose connectome have been mapped [1]. There are two types of connections between neurons: chemical(chemical synapses) and electrical(junction potentials).

INSERT REASON for 279 vertex graph instead of 302.

The objective of the experiment is to match the chemical graph G_1 to the electrical one G_2 . The adjacency matrices for both graphs are sparse. Both G_1 and G_2 are weighted graphs, i.e each edge has a nonzero real number attribute, so for sake of uniformity with the other examples, the adjacency matrices are binarized and symmetrized. For each number of hard seeds, $m \in \{0, 1, 5, 10, 20, 50, 75, 100, 150, 200\}$, there are different number of vertices to be matched $n = k - m$. Since the expected number of matches is always 1, the match ratio for random chance $\frac{1}{n}$ is different for each different number of hard seeds. Note the increasing random chance curve in Fig. 3.

Although the match ratio is generally low which should be expected due to the difficulty of the problem, increasing number of hard seeds results in a consistent increase in match ratio.

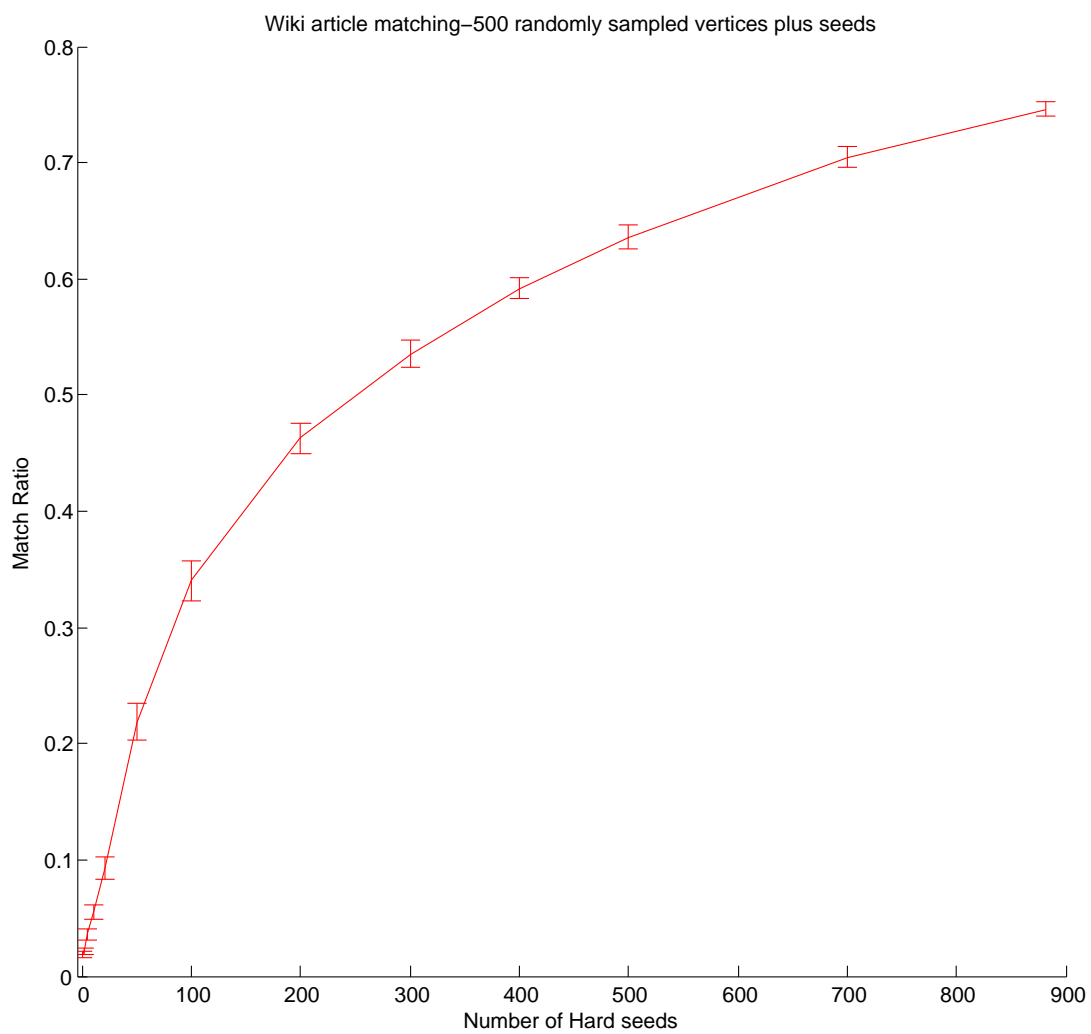


Figure 1: Matching English and French subgraphs




Figure 2: Matching English and French subgraphs

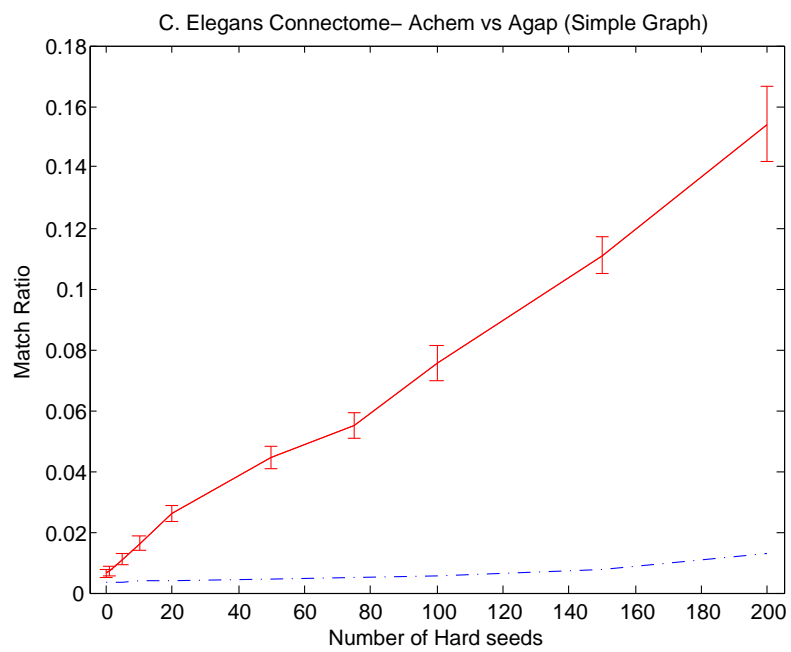


Figure 3: Matching of Chemical and Electrical connectivity graphs of *C. elegans* nervous system

8 Enron example

Enron email corpus consists of mails between $k = 184$ employees in Enron. Publicly available emails are used to compute a time-series of graphs $\{G_t, t = 1, \dots, T\}$ between k vertices. The inference task is to find anomalies (schemers scheming to commit fraud) in this communication graph. Previous work [1] indicates suspicious behaviour takes place at time stamps $t = (130, 131, 132)$. Graphs with consecutive timesteps at $t = (130, 131, 132)$ are matched one pair at a time using modified-FAQ. The match ratio for each pair is plotted in 4.

The results are consistent with previous work [1], the average match ratio is much higher between the graphs at $t = 130$ and $t = 131$ compared to matching between $t = 131$ and $t = 132$ and between $t = 130$ and $t = 132$. The anomalous event in the graph time-series takes place at $t = 132$ according to [1], which makes the graph matching more difficult compared to matching graphs at other times.

Note that the Two-Class experiment is an appropriate model for this problem. The list of anomalous vertices (scheming fraudsters), as detected by [1], are $\{v \in V : \gamma(v) = 1\} = \{2, 3, 4, 5, 7, 10, 11, 17, 18, 19, 22, 28, 31, 33, 40, 41, 44, 46, 47, 48, 50, 52, 54, 55, 57, 59, 60, 63, 64, 65, 71, 76, 77, 79, 80, 83, 84, 85, 86, 87, 89, 90, 94, 97, 98, 102, 103, 107, 108, 113, 115, 117, 119, 120, 126, 127, 128, 130, 133, 134, 137, 138, 140, 141, 144, 146, 147, 148, 149, 150, 155, 157, 158, 160, 164, 165, 167, 174, 176, 177, 180, 183\}$

References

- [1] H.A. Almohamad and S.O. Duffuaa, A linear programming approach for the weighted graph matching problem, *IEEE Transactions on Pattern Analysis and Machine Intelligence* **15** (1993) pp 522–525.
- [2] D. Conte, P. Foggia, C. Sansone, and M. Vento, Thirty years of graph matching in pattern recognition, *International Journal of Pattern Recognition and Artificial Intelligence* **18** (2004) pp 265–298.
- [3] J.T. Vogelstein, J.M. Conroy, L.J. Podrazik, S.G. Kratzer, E.T. Harley, D.E. Fishkind, R.J. Vogelstein, and C.E. Priebe, Brain graph matching via fast approximate quadratic programming, submitted. Available at arxiv.org/pdf/1112.5507

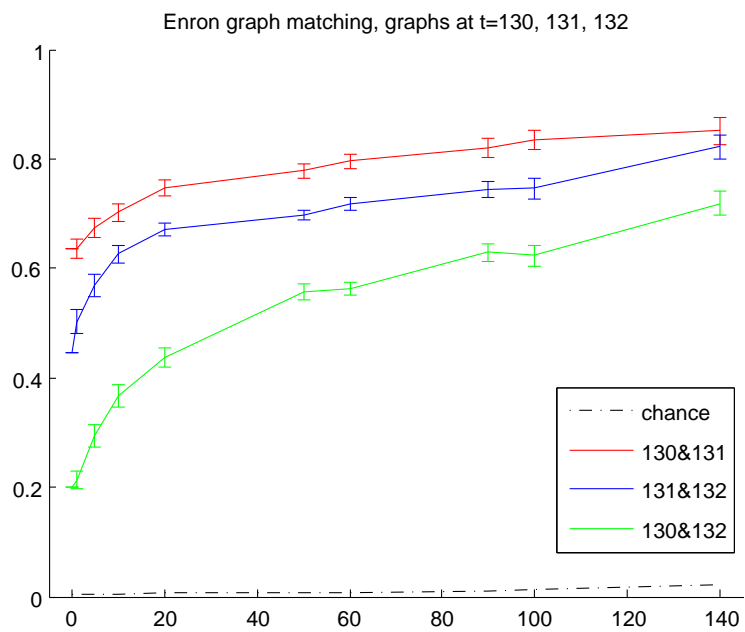


Figure 4: Matching of Enron communication graph at timestamps: 130,131,132